

CSC 3210 Computer organization and programming

Chapter 3

Chunlan Gao



Review



IEEE Std 754-1985

Bias: Single precision 127, double precision 1203

1111 1110 ---> 00000001 ?

↓
-126 + 127 --> 1, so the bias is 127

normal value :

Exponents from 00000001 to 11111110

$$(-1)^S \times (1 + \text{Fraction}) \times 2^E$$

Reserved:

Exponents 00000000 and 11111111

Denormal Numbers (Subnormal numbers)

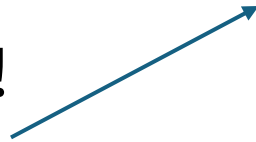
- Exponent = 000...0 \Rightarrow hidden bit is 0

$$x = (-1)^s \times (0 + \text{Fraction}) \times 2^{-\text{Bias}}$$

- Smaller than normal numbers (at $\approx 1.4 \times 10^{-45}$ and go down to zero)
 - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0

$$x = (-1)^s \times (0 + 0) \times 2^{-\text{Bias}} = \pm 0.0$$

Two representations of 0.0!



Infinites and NaNs



- Exponent = 111...1, Fraction = 000...0
 - \pm Infinity
 - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction \neq 000...0
 - Not-a-Number (NaN)
 - Indicates illegal or undefined result
 - e.g., $0.0 / 0.0$
 - Can be used in subsequent calculations

Floating-Point Addition



- Consider a 4-digit decimal example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2

Floating-Point Addition



Now consider a 4-digit binary example

$0.5 + 0.4375$

$$0.5_{\text{ten}} = 1/2 = 0.1_{\text{two}} = 1.000 \times 2^{-1}$$

$$0.4375 = -7/16 = -7/2^4 = -0.0111_{\text{two}} = -1.110_{\text{two}} \times 2^{-2}$$

Step1:

The significand of the number with the lesser exponent ($-1.11_{\text{two}} \times 2^{-2}$) is shifted right until its exponent matches the larger number

$$-1.110_{\text{two}} \times 2^{-2} = -0.111_{\text{two}} \times 2^{-1}$$

Step 2. Add the significands: $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$

- 3. Normalize result & check for over/underflow: Since $127 \geq -4 \geq -126$, there is no overflow or underflow.

$$\begin{aligned} 0.001_{\text{two}} \times 2^{-1} &= 0.010_{\text{two}} \times 2^{-2} = 0.100_{\text{two}} \times 2^{-3} \\ &= 1.000_{\text{two}} \times 2^{-4} \end{aligned}$$

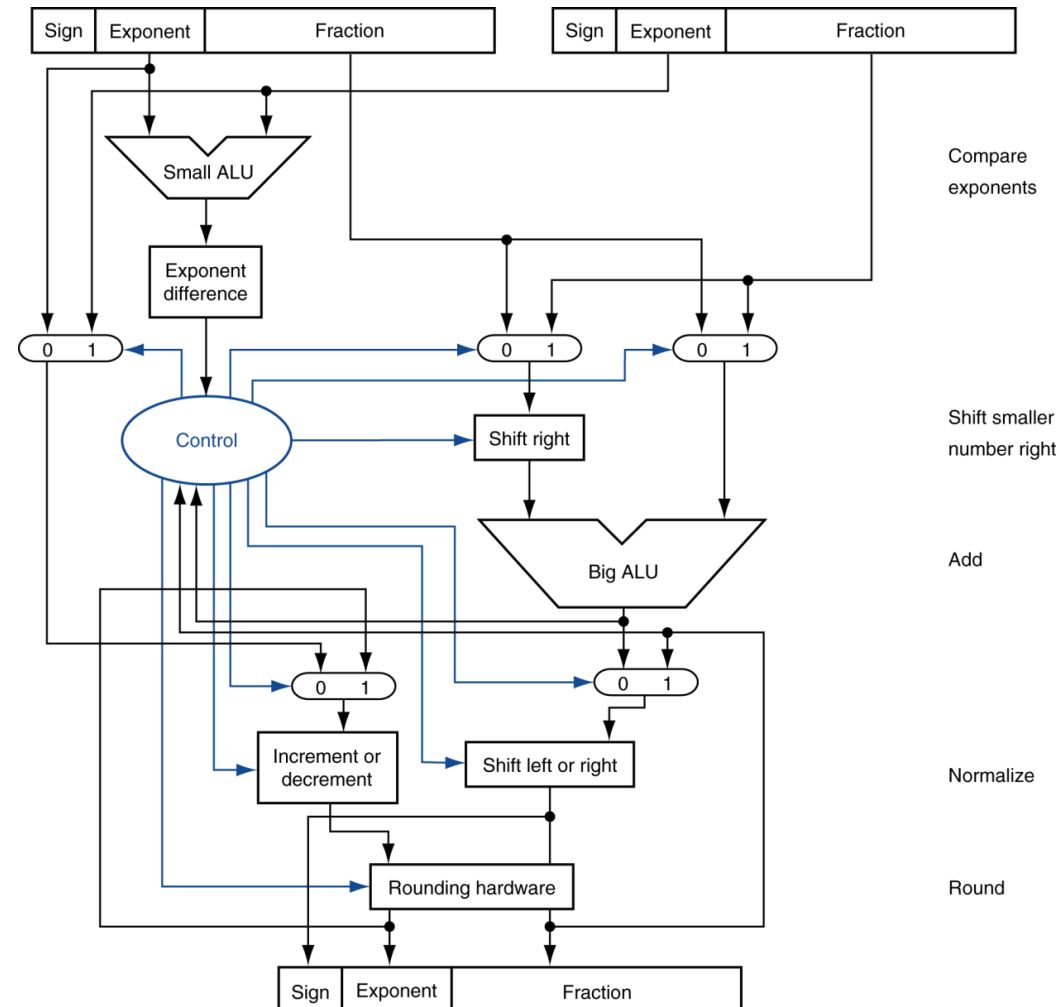
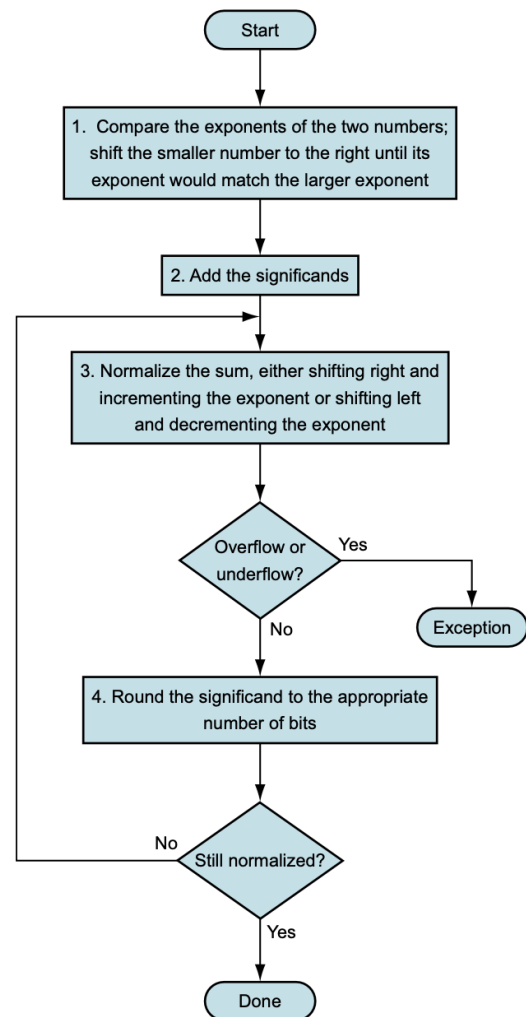
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

FP Adder Hardware



- Much more complex than integer adder
- Doing it in one clock cycle would take too long
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- FP adder usually takes several cycles
 - Can be pipelined

FP Adder Hardware



Floating-Point Multiplication



- Consider a 4-digit decimal example
 - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
 - For biased exponents, subtract bias from sum
 - New exponent = $10 + -5 = 5$
- 2. Multiply significands
 - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- 3. Normalize result & check for over/underflow
 - 1.0212×10^6
- 4. Round and renormalize if necessary
 - 1.021×10^6
- 5. Determine sign of result from signs of operands
 - $+1.021 \times 10^6$

Floating-Point Multiplication



- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$ (0.5×-0.4375)
- 1. Add exponents
 - Unbiased: $-1 + -2 = -3$
 - Biased: $(-1 + 127) + (-2 + 127) = -3 + 254 - 127 = -3 + 127$
- 2. Multiply significands
 - $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
 - $1.110_2 \times 2^{-3}$ (no change) with no over/underflow
- 4. Round and renormalize if necessary
 - $1.110_2 \times 2^{-3}$ (no change)
- 5. Determine sign: $+ve \times -ve \Rightarrow -ve$
 - $-1.110_2 \times 2^{-3} = -0.21875$

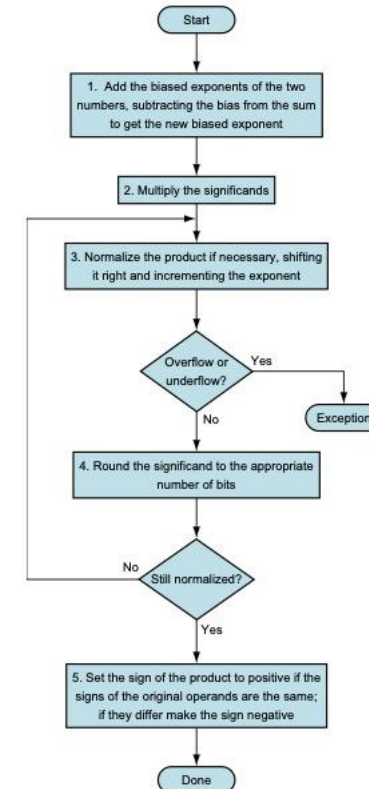


FIGURE 3.16 Floating-point multiplication. The normal path is to execute steps 3 and 4 once, but if rounding causes the sum to be unnormalized, we must repeat step 3.

FP Arithmetic Hardware



- FP multiplier is of similar complexity to FP adder
 - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
 - Addition, subtraction, multiplication, division, reciprocal, square-root
 - $\text{FP} \leftrightarrow \text{integer}$ conversion
- Operations usually takes several cycles
 - Can be pipelined

FP Instructions in RISC-V



- Separate FP registers: f0, ..., f31
 - double-precision
 - single-precision values stored in the lower 32 bits
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - flw, fld
 - fsw, fsd

FP Instructions in RISC-V



- Single-precision arithmetic
 - `fadd.s`, `fsub.s`, `fmul.s`, `fdiv.s`, `fsqrt.s`
 - e.g., `fadds.s f2, f4, f6`
- Double-precision arithmetic
 - `fadd.d`, `fsub.d`, `fmul.d`, `fdiv.d`, `fsqrt.d`
 - e.g., `fadd.d f2, f4, f6`
- Single- and double-precision comparison
 - `feq.s`, `flt.s`, `fle.s`
 - `feq.d`, `flt.d`, `fle.d`
 - Result is 0 or 1 in integer destination register
 - Use `beq`, `bne` to branch on comparison result
- Branch on FP condition code true or false
 - `B.cond`

FL in RISC-V



RISC-V floating-point assembly language

Arithmetic	FP add single	<code>fadd.s f0, f1, f2</code>	$f0 = f1 + f2$	FP add (single precision)
	FP subtract single	<code>fsub.s f0, f1, f2</code>	$f0 = f1 - f2$	FP subtract (single precision)
	FP multiply single	<code>fmul.s f0, f1, f2</code>	$f0 = f1 * f2$	FP multiply (single precision)
	FP divide single	<code>fdiv.s f0, f1, f2</code>	$f0 = f1 / f2$	FP divide (single precision)
	FP square root single	<code>fsqrt.s f0, f1</code>	$f0 = \sqrt{f1}$	FP square root (single precision)
	FP add double	<code>fadd.d f0, f1, f2</code>	$f0 = f1 + f2$	FP add (double precision)
	FP subtract double	<code>fsub.d f0, f1, f2</code>	$f0 = f1 - f2$	FP subtract (double precision)
	FP multiply double	<code>fmul.d f0, f1, f2</code>	$f0 = f1 * f2$	FP multiply (double precision)
	FP divide double	<code>fdiv.d f0, f1, f2</code>	$f0 = f1 / f2$	FP divide (double precision)
	FP square root double	<code>fsqrt.d f0, f1</code>	$f0 = \sqrt{f1}$	FP square root (double precision)
Comparison	FP equality single	<code>feq.s x5, f0, f1</code>	$x5 = 1$ if $f0 == f1$, else 0	FP comparison (single precision)
	FP less than single	<code>flt.s x5, f0, f1</code>	$x5 = 1$ if $f0 < f1$, else 0	FP comparison (single precision)
	FP less than or equals single	<code>fle.s x5, f0, f1</code>	$x5 = 1$ if $f0 \leq f1$, else 0	FP comparison (single precision)
	FP equality double	<code>feq.d x5, f0, f1</code>	$x5 = 1$ if $f0 == f1$, else 0	FP comparison (double precision)
	FP less than double	<code>flt.d x5, f0, f1</code>	$x5 = 1$ if $f0 < f1$, else 0	FP comparison (double precision)
	FP less than or equals double	<code>fle.d x5, f0, f1</code>	$x5 = 1$ if $f0 \leq f1$, else 0	FP comparison (double precision)
Data transfer	FP load word	<code>flw f0, 4(x5)</code>	$f0 = \text{Memory}[x5 + 4]$	Load single-precision from memory
	FP load doubleword	<code>fld f0, 8(x5)</code>	$f0 = \text{Memory}[x5 + 8]$	Load double-precision from memory
	FP store word	<code>fsw f0, 4(x5)</code>	$\text{Memory}[x5 + 4] = f0$	Store single-precision from memory
	FP store doubleword	<code>fsd f0, 8(x5)</code>	$\text{Memory}[x5 + 8] = f0$	Store double-precision from memory

FIGURE 3.17 RISC-V floating-point architecture revealed thus far. This information is also found in column 2 of the RISC-V Reference Data Card at the front of this book.

FP Example: °F to °C



- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0f));  
}
```

- fahr in f10, result in f10, literals in global memory space
- We assume that the compiler places the three floating-point constants in memory within easy reach of register x3

- Compiled RISC-V code:

f2c:

```
f1w    f0,const5(x3)  // f0 = 5.0f  
f1w    f1,const9(x3)  // f1 = 9.0f  
fdiv.s f0, f0, f1     // f0 = 5.0f / 9.0f  
f1w    f1,const32(x3) // f1 = 32.0f  
fsub.s f10,f10,f1     // f10 = fahr - 32.0  
fmul.s f10,f0,f10     // f10 = (5.0f/9.0f) * (fahr-32.0f)  
jalr   x0,0(x1)       // return
```

$$^{\circ}C = \frac{5}{9} \times (^{\circ}F - 32)$$

FP Example: Array Multiplication



- $C = C + A \times B$
 - All 32×32 matrices, 64-bit double-precision elements

- C code:

```
void mm (double c[][],
         double a[][], double b[][]) {
    size_t i, j, k;
    for (i = 0; i < 32; i = i + 1)
        for (j = 0; j < 32; j = j + 1)
            for (k = 0; k < 32; k = k + 1)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
}
```

- Addresses of c, a, b in x10, x11, x12, and i, j, k in x5, x6, x7
- Assume that the integer variables are in x5, x6, and x7, respectively

Accurate Arithmetic



Index	Element	Address offset (bytes)
0	A[0][0]	0
1	A[0][1]	8
2	A[0][2]	16
...
31	A[0][31]	248
32	A[1][0]	256
33	A[1][1]	264
...
1023	A[31][31]	8184



FP Example: Array Multiplication



- RISC-V code:

mm: ...

```
li    x28,32      // x28 = 32 (row size/loop end)
li    x5,0        // i = 0; initialize 1st for loop
```

```
L1:   li    x6,0    // j = 0; initialize 2nd for loop
```

```
L2:   li    x7,0    // k = 0; initialize 3rd for loop
```

f0 = c[i][j]

```
slli  x30,x5,5     // x30 = i * 2**5 (size of row of c)
add   x30,x30,x6    // x30 = i * size(row) + j
slli  x30,x30,3     // x30 = byte offset of [i][j]
add   x30,x10,x30   // x30 = byte address of c[i][j]
fld   f0,0(x30)    // f0 = c[i][j]
```

f1 = b[k][j] and f2= a[k][j]

```
L3:   slli  x29,x7,5 // x29 = k * 2**5 (size of row of b)
add   x29,x29,x6    // x29 = k * size(row) + j
slli  x29,x29,3     // x29 = byte offset of [k][j]
add   x29,x12,x29   // x29 = byte address of b[k][j]
fld   f1,0(x29)    // f1 = b[k][j]
```

FP Example: Array Multiplication



```
slli    x29,x5,3      // x29 = i * 2**5 (size of row of a)
add     x29,x29,x7     // x29 = i * size(row) + k
slli    x29,x29,3      // x29 = byte offset of b[i][k]
add     x29,x11,x29    // x29 = byte address of a[i][k]
fld     f2,0(x29)      // f2 = a[i][k]
#f[0]= c[i][j] + a[i][k] * b[k][j]
fmul.d  f1, f2, f1     // f1 = a[i][k] * b[k][j]
fadd.d  f0, f0, f1     // f0 = c[i][j] + a[i][k] * b[k][j]
# k++, check if k<32
addi    x7,x7,1        // k = k + 1
bltu    x7,x28,L3      // if (k < 32) go to L3
# save f0 to c[i][j]
fsd     f0,0(x30)      // c[i][j] = f0
addi    x6,x6,1        // j = j + 1
bltu    x6,x28,L2      // if (j < 32) go to L2
addi    x5,x5,1        // i = i + 1
bltu    x5,x28,L1      // if (i < 32) go to L1
```

Accurate Arithmetic



IEEE 754, therefore, always keeps two extra bits on the right during intervening additions, called **guard** and **round**

Add $2.56 * 10^0$ to $2.34 * 10^2$, assuming we have three significant decimal digit, (the guard digit hold 5 and round digit hold 6,)

$$\begin{array}{r} 2.3400_{\text{ten}} \\ + 0.0256_{\text{ten}} \\ \hline 2.3656_{\text{ten}} \end{array}$$

Round to 3 digit 2.37

Without guard and round
the result is 2.36

$$\begin{array}{r} 2.34_{\text{ten}} \\ + 0.02_{\text{ten}} \\ \hline 2.36_{\text{ten}} \end{array}$$

Subword Parallelism

